

# Using Elasticsearch to Help Monitor Docker Containers

---

OCTOBER 19, 2016



# Tonight's Discussion

---

- Introduction into the Elastic Stack and Beats
- Overview of tonight's demonstration setup
- Topbeat configuration
- Assembling the Elastic Stack
- Assembling the sample application
- Interspersed with Docker 1.12 Swarm and Routing Mesh



# Tonight's Discussion

---

- We'll be moving pretty quickly
  - This presentation is targeted toward an intermediate level
    - Some a bit more advanced level
  - Due to time, not much looking into the Dockerfiles
    - Happy to discuss after the presentation
    - They will be made available
- Do ask questions
  - Some may be deferred to the end of the session
  - There will be a Q&A at the end



# All Artifacts will be Available on GitHub

*In few days; maybe early next week.*

*A note will be posted on the Meetup site as soon as they are available.*



# The Elastic Stack

---

# The Elastic Stack

---



## Elasticsearch 2.4.1

a distributed, open source search and analytics engine



## Kibana 4.6.1

an open source data visualization platform that allows you to interact with Elasticsearch data



## Beats 1.3.1

open source data shippers for Elasticsearch



## Marvel 2.4.0

provides status for an Elasticsearch deployment



Soon to be version 5.0 across the suite



# Shipping Data to Elasticsearch

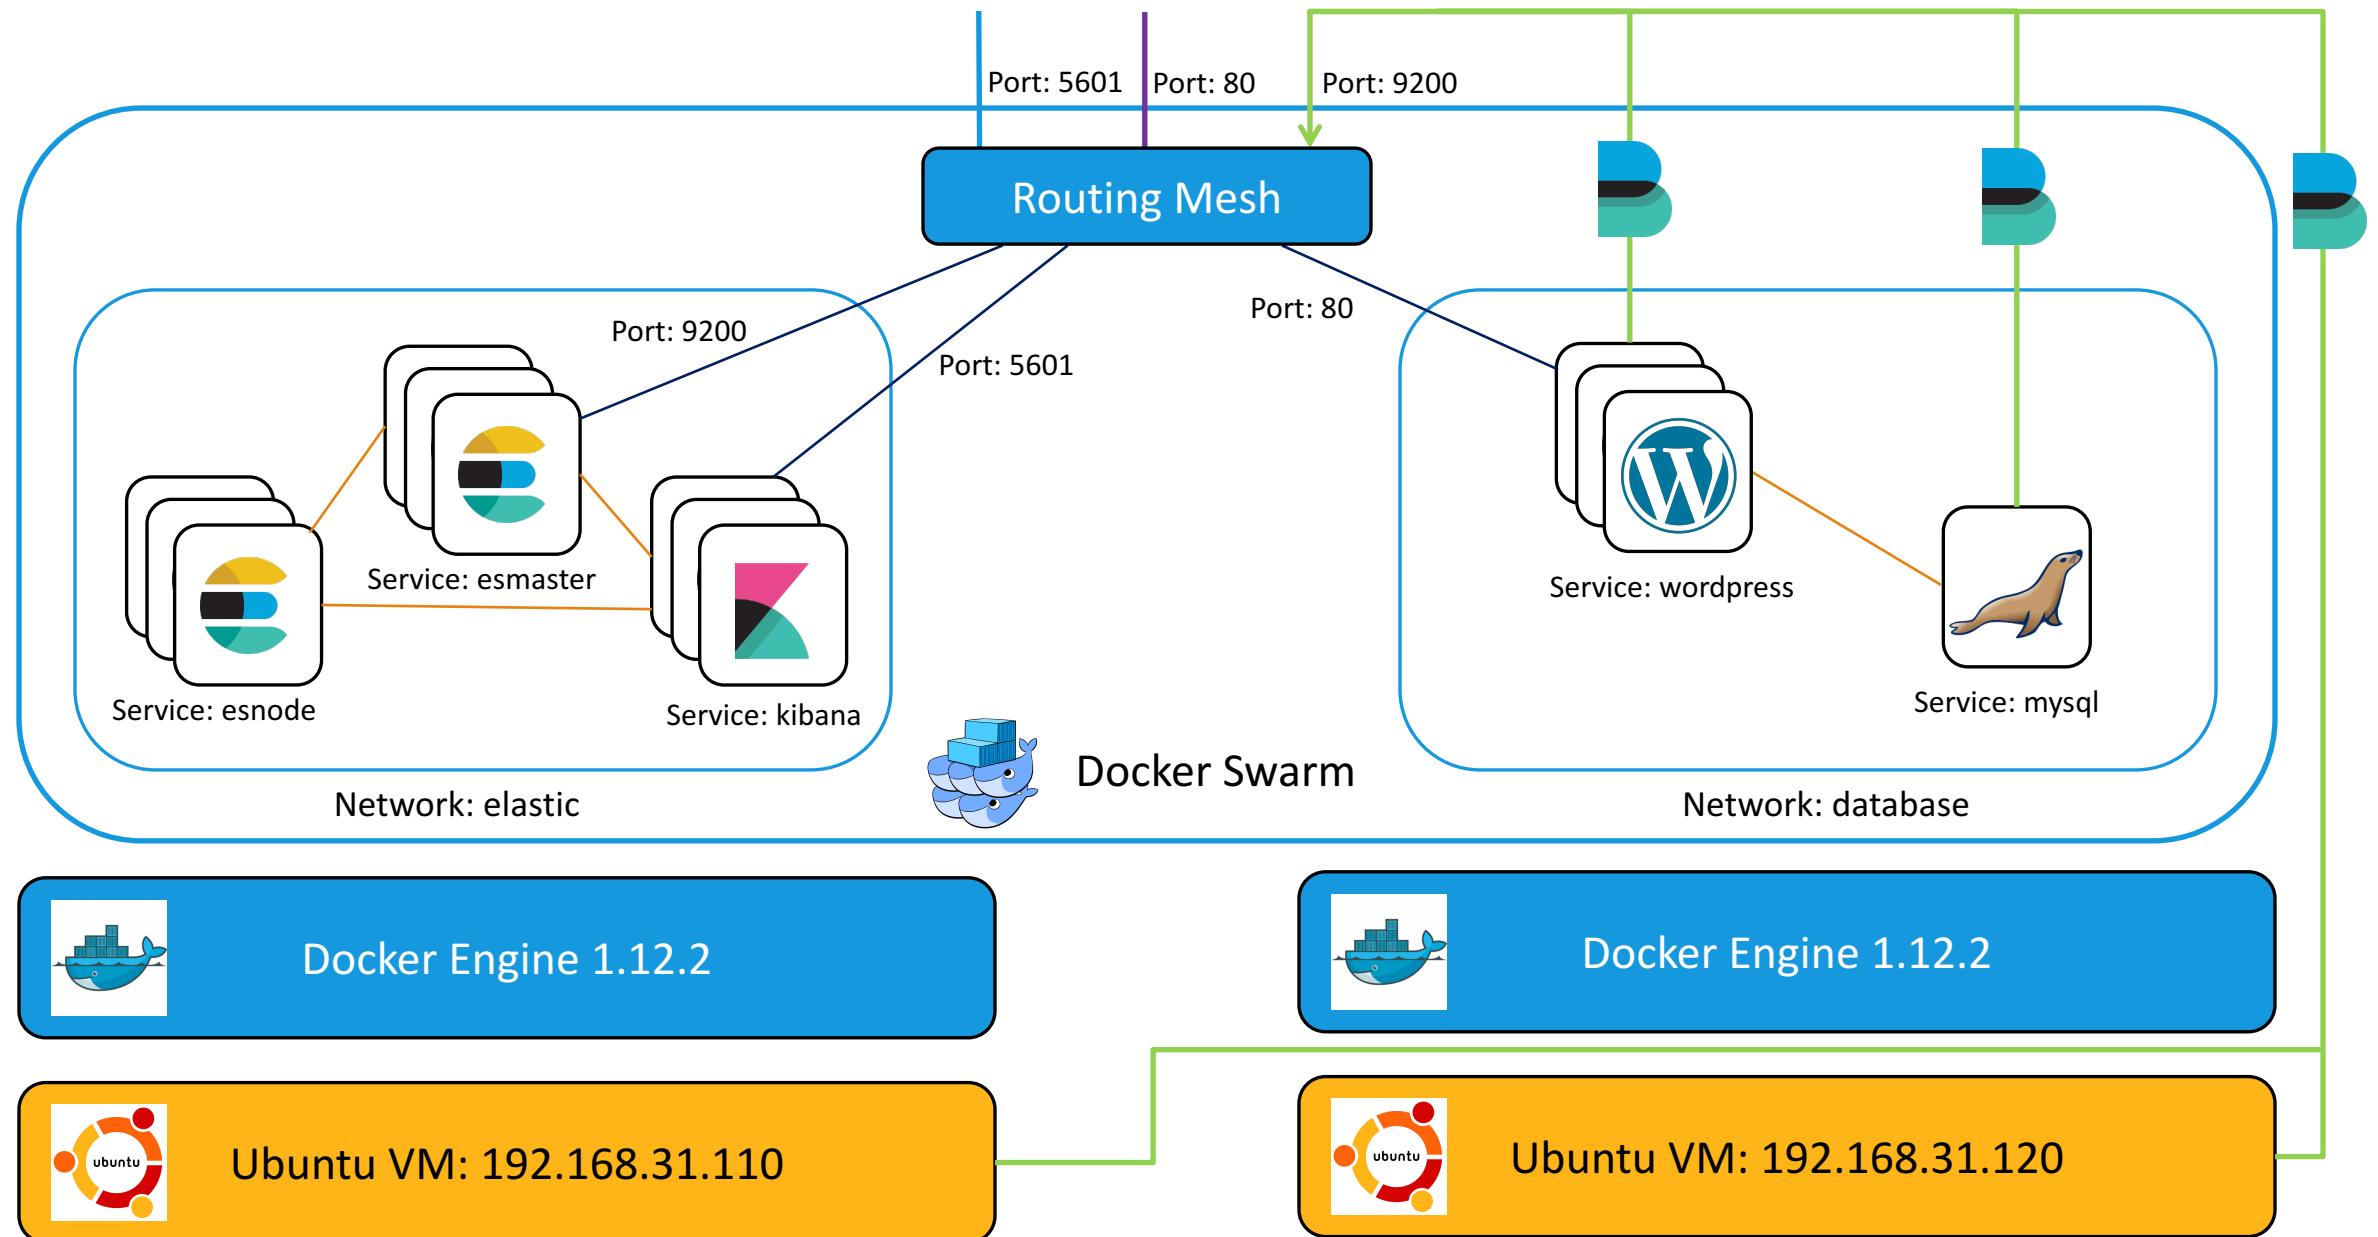
---

- Filebeat: Analyze Log Files in Real Time
- Packetbeat: Tap into Your Wire Data
- Winlogbeat: Gather Insight from Windows Event Logs
- Topbeat: Gather Infrastructure Metrics
- Libbeat: For Building More Beats
- Several community contributions as well



# Demonstration Setup

---



# Topbeat Configuration

---

# Capturing System and Process Data

---

- System-wide
  - CPU, memory, swap, etc.
- Per-process
  - Name, PID, state, CPU, memory
- File system
  - Disk mount information, available space



# Straightforward Configuration

---

## topbeat.yml

```
input:  
  # In seconds, defines how often to read server statistics  
  period: 10  
  
  # Regular expression to match the processes that are monitored  
  # By default, all the processes are monitored  
  procs: [".*"]  
  
  # Statistics to collect (all enabled by default)  
  stats:  
    # per system statistics, by default is true  
    system: true  
  
    # per process statistics, by default is true  
    process: true  
  
    # file system information, by default is true  
    filesystem: true  
  
    # cpu usage per core, by default is false  
    cpu_per_core: false
```



# Two Configurations for the Demo

---

- Topbeat on the VMs
  - VM system-wide data
  - Without per-process data
- Topbeat in the containers
  - Per-process data
  - Without system-wide data

```
# per system statistics, by default is true
system: true

# per process statistics, by default is true
process: false
```

```
# per system statistics, by default is true
system: false

# per process statistics, by default is true
process: true
```

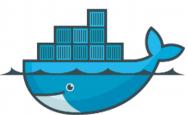
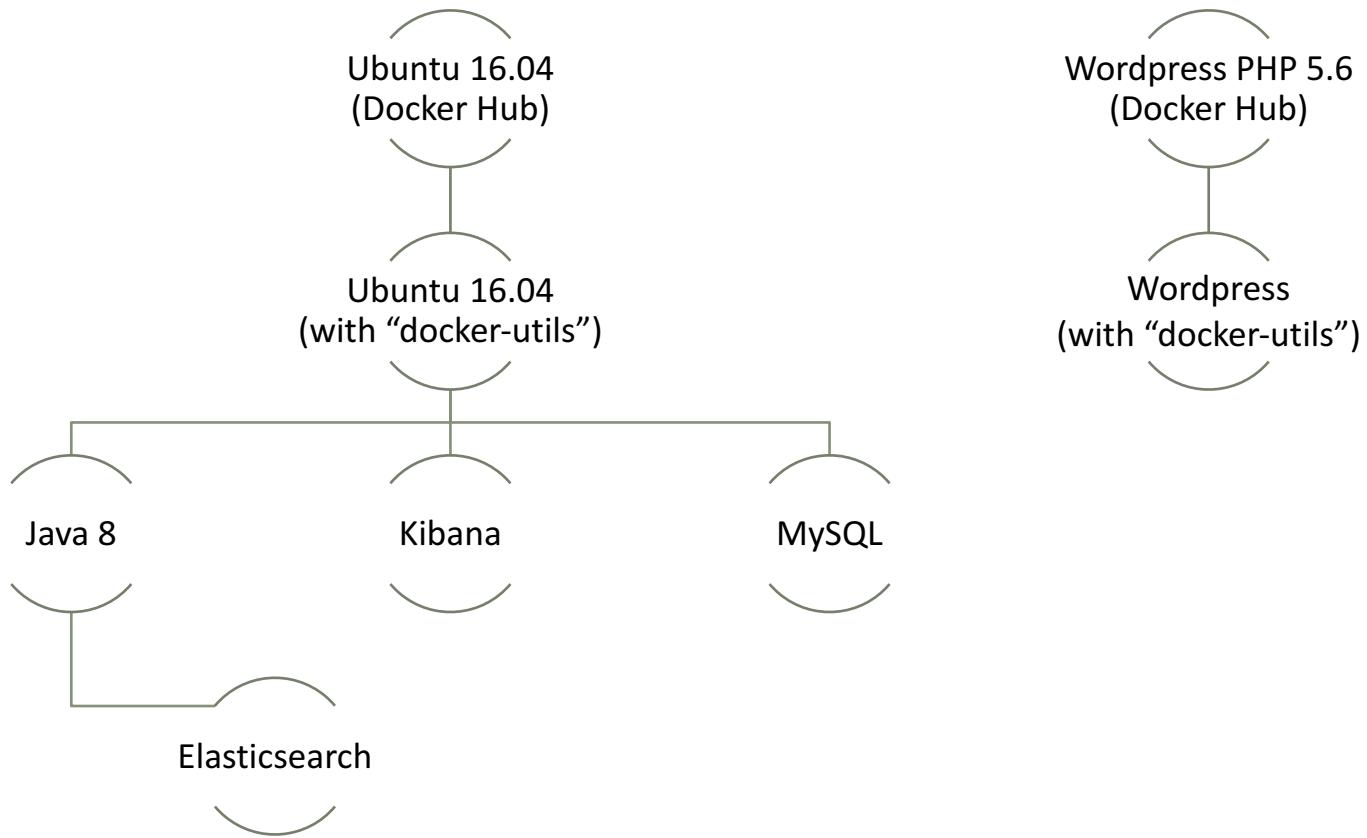


# Image Hierarchy

---

# Image Hierarchy

---



# Assembling the Elastic Stack

---

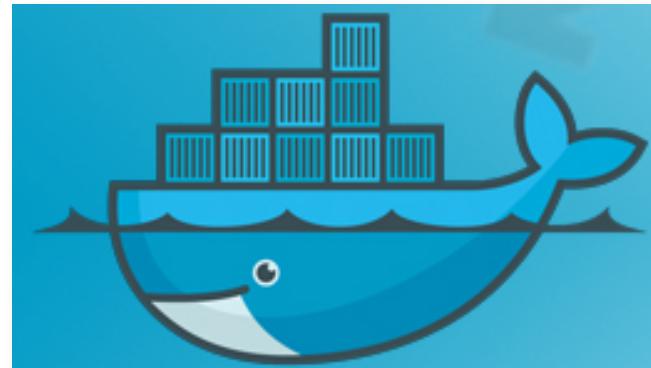
# Elasticsearch Master Node Service

---

```
docker service create \
--replicas 3 \
--publish 9200:9200 \
--name esmaster \
--network elastic \
--mount type=bind,source=/opt/esdata,destination=/home/elasticsearch/data \
meetup/elasticsearch:2.4.1 \
/home/elasticsearch/bin/start-elasticsearch.sh -m 2
```



# Let's Take a Look



**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

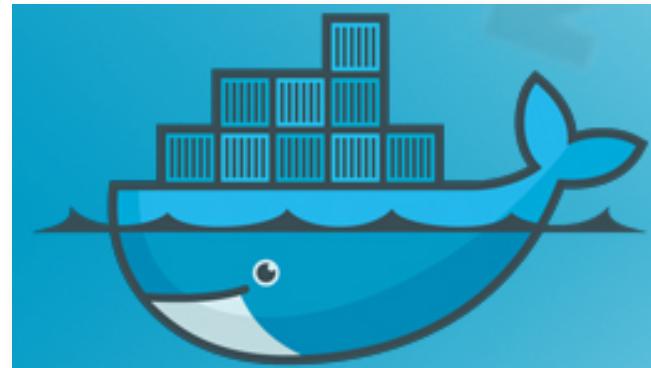
# Kibana Service

---

```
docker service create \
  --publish 5601:5601 \
  --name kibana \
  --network elastic \
  meetup/kibana:4.6.1 \
  /home/kibana/bin/start-kibana.sh -s esmaster
```



# Let's Take a Look

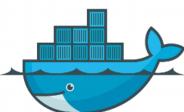


**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

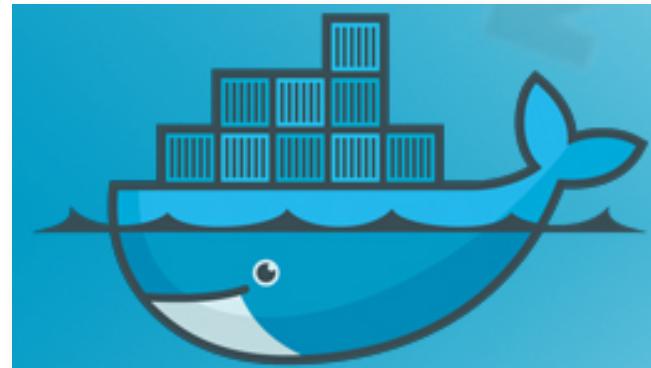
# Elasticsearch Worker Node Service

---

```
docker service create \
--replicas 2 \
--name esnode \
--network elastic \
--mount type=bind,source=/opt/esdata,destination=/home/elasticsearch/data \
meetup/elasticsearch:2.4.1 \
/home/elasticsearch/bin/start-elasticsearch.sh -m 2
```



# Let's Take a Look



**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

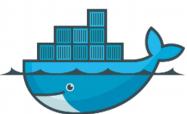
# Assembling the Sample Application

---

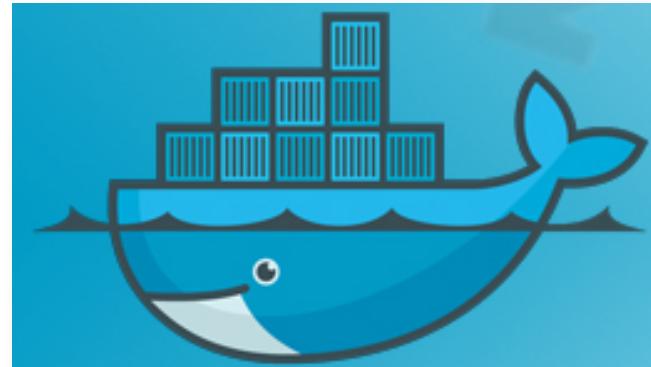
# MySQL Service

---

```
docker service create \
--name mysql \
--network database \
--env ES_SERVER=192.168.31.110 \
--env MYSQL_ROOT_PASSWORD=mysqlroot \
--constraint engine.labels.meetup.node==mysql \
meetup/mysql:5.7
```



# Let's Take a Look

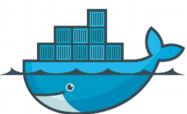


**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

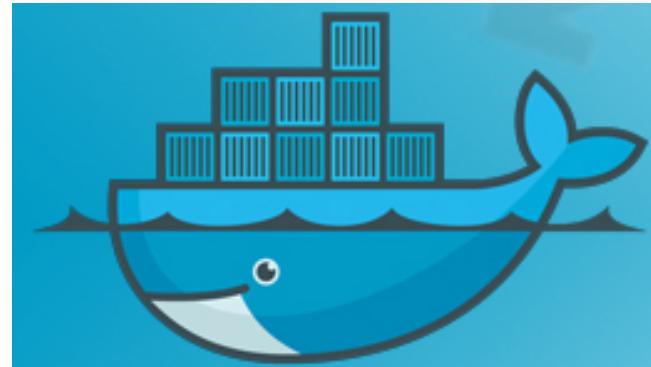
# Wordpress Service

---

```
docker service create \
--name wordpress \
--network database \
--publish 80:80 \
--env ES_SERVER=192.168.31.110 \
--env WORDPRESS_DB_PASSWORD=mysqlroot \
--mount type=bind,source=/opt/wordpress,destination=/var/www/html \
--constraint engine.labels.meetup.node==wordpress \
meetup/wordpress:php5.6
```

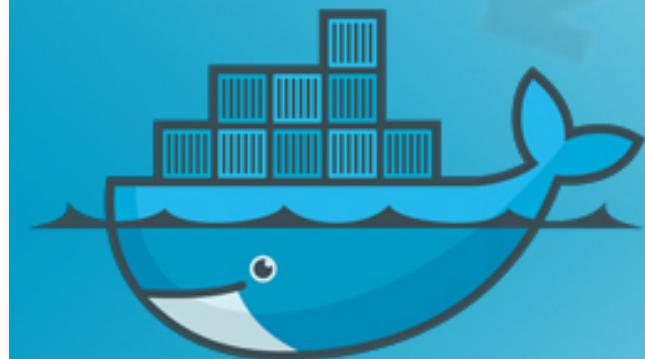


# Let's Take a Look



**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

# Questions?



**DOCKER ATLANTA**  
Build, Ship and Run Any App, Anywhere  
**#dockermeetup**

# Using Elasticsearch to Help Monitor Docker Containers

---

OCTOBER 19, 2016

